

Note on CHIME (not CHORD) FRB beamforming

Kendrick Smith

July 15, 2024

1 First FFT (256 \rightarrow 512)

We want to compute the following FFT:

$$Y_q = \sum_n \underbrace{e^{2\pi i q n / 512}}_{\equiv E(q, n)} X_n \quad 0 \leq q < 512 \quad 0 \leq n < 256 \quad (1)$$

The output index $0 \leq q < 512$ has 9 bits, and the input index $0 \leq n < 256$ has 8 bits:

$$q = q_8 q_7 q_6 q_5 q_4 q_3 q_2 q_1 q_0 \quad n = n_7 n_6 n_5 n_4 n_3 n_2 n_1 n_0 \quad (2)$$

In this section, we'll present an in-register "microkernel" with the following register assignments for the input/output arrays:

$$[X] \quad b \leftrightarrow \text{ReIm} \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow n_7 n_6 n_5 n_4 n_3 \quad r \leftrightarrow n_2 n_1 n_0 \quad (3)$$

$$[Y] \quad b \leftrightarrow \text{ReIm} \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow q_4 q_5 q_6 q_7 q_8 \quad r \leftrightarrow q_0 q_1 q_2 q_3 \quad (4)$$

1.1 Overview

We can write the phase $E(q, n)$ defined in (1) in terms of index bits as follows:

$$\begin{aligned} E(q, n) &= \exp \left[\frac{2\pi i}{512} \left(\sum_j q_j 2^j \right) \left(\sum_k n_k 2^k \right) \right] \\ &= \prod_{jk} \underbrace{\exp \left[\frac{2\pi i}{512} q_j n_k 2^{j+k} \right]}_{\equiv E(q_j, n_k)} \end{aligned} \quad (5)$$

Note that $E(q_j, n_k) = 1$ if $j + k \geq 9$.

Now we define the following generalization of the $E(q_j, n_k)$ notation, which is easiest to define by example. Consider a symbol like $E(q_0 q_1 q_2, n_0 n_1 n_6 n_7)$ whose first argument is a subset of the q -bits, and whose second argument is a subset of the n -bits. We define it as follows:

$$E(q_0 q_1 q_2, n_0 n_1 n_6 n_7) \equiv \prod_{\substack{j \in \{0,1,2\} \\ k \in \{0,1,6,7\}}} E(q_j, n_k) \quad (6)$$

Our FFT microkernel is based on the following factorization of the phase $E(q, n)$ appearing in Eq. (1):

$$\begin{aligned} E(q, n) &= E(q_6 q_7 q_8, n_0 n_1 n_2) E(q_3 q_4 q_5, n_0 n_1 n_2) E(q_3 q_4 q_5, n_3 n_4 n_5) \\ &\quad \times E(q_0 q_1 q_2, n_2 n_3 n_4 n_5) E(q_0 q_1 q_2, n_0 n_1 n_6 n_7) \end{aligned} \quad (7)$$

leading to the following sequence ($X \rightarrow Z_1 \rightarrow Z_2 \rightarrow Z_3 \rightarrow Z_4 \rightarrow Y$) of array operations:

$$Z_1[q_0q_1q_2n_0n_1n_2n_3n_4n_5] = \sum_{n_6n_7} E(q_0q_1q_2, n_0n_1n_6n_7)X_n \quad (8)$$

$$Z_2[q_0q_1q_2n_0n_1n_2n_3n_4n_5] = E(q_0q_1q_2, n_2n_3n_4n_5)Z_1[q_0q_1q_2n_0n_1n_2n_3n_4n_5] \quad (\text{no sum}) \quad (9)$$

$$Z_3[q_0q_1q_2q_3q_4q_5n_0n_1n_2] = \sum_{n_3n_4n_5} E(q_3q_4q_5, n_3n_4n_5)Z_2[q_0q_1q_2n_0n_1n_2n_3n_4n_5] \quad (10)$$

$$Z_4[q_0q_1q_2q_3q_4q_5n_0n_1n_2] = E(q_3q_4q_5, n_0n_1n_2)Z_3[q_0q_1q_2q_3q_4q_5n_0n_1n_2] \quad (\text{no sum}) \quad (11)$$

$$Y_q = \sum_{n_0n_1n_2} E(q_6q_7q_8, n_0n_1n_2)Z_4[q_0q_1q_2q_3q_4q_5n_0n_1n_2] \quad (12)$$

These computational steps can be done as follows:

- To compute $X \rightarrow Z_1$ using Eq. (8), we use **m16n8k8** MMAs (Appendix B.1).
- To compute $Z_1 \rightarrow Z_2$ using Eq. (9), we use **__half2** FMAs.
- To compute $Z_2 \rightarrow Z_3$ using Eq. (10), we use **m16n8k16** MMAs (Appendix B.2).
- To compute $Z_3 \rightarrow Z_4$ using Eq. (11), we use **__half2** FMAs.
- To compute $Z_4 \rightarrow Y$ using Eq. (12), we use **m16n8k16** MMAs (Appendix B.2).

1.2 Register assignments

The register assignments for the X, Y, Z arrays are:

$$[X] \quad b \leftrightarrow \text{ReIm} \quad t_0t_1t_2t_3t_4 \leftrightarrow n_7n_6n_5n_4n_3 \quad r \leftrightarrow n_2n_1n_0 \quad (13)$$

$$[Z_1, Z_2] \quad b \leftrightarrow n_5 \quad t_0t_1t_2t_3t_4 \leftrightarrow n_4n_3q_0q_1q_2 \quad r \leftrightarrow n_2n_1n_0\text{ReIm} \quad (14)$$

$$[Z'_3] \quad b \leftrightarrow q_0 \quad t_0t_1t_2t_3t_4 \leftrightarrow q_1q_2q_3q_4q_5 \quad r \leftrightarrow n_2n_1n_0\text{ReIm} \quad (15)$$

$$[Z_3, Z_4] \quad b \leftrightarrow n_2 \quad t_0t_1t_2t_3t_4 \leftrightarrow n_1n_0q_3q_4q_5 \quad r \leftrightarrow q_0q_1q_2\text{ReIm} \quad (16)$$

$$[Y'] \quad b \leftrightarrow q_3 \quad t_0t_1t_2t_3t_4 \leftrightarrow q_4q_5q_6q_7q_8 \quad r \leftrightarrow q_0q_1q_2\text{ReIm} \quad (17)$$

$$[Y] \quad b \leftrightarrow \text{ReIm} \quad t_0t_1t_2t_3t_4 \leftrightarrow q_4q_5q_6q_7q_8 \quad r \leftrightarrow q_0q_1q_2q_3 \quad (18)$$

Note that for Z_3 and Y , we have given two register assignments (“primed” and “unprimed”). In both cases, we go from the primed assignment to the unprimed assignment using local/warp transpose operations (Appendix A).

For the main computational steps ($X \rightarrow Z_1 \rightarrow Z_2 \rightarrow Z'_3$) and ($Z_3 \rightarrow Z_4 \rightarrow Y'$), we need some precomputed phases in registers. We’ll use the following notation for these arrays of precomputed phases (see Eqs. (8)–(12)):

$$E_{X \rightarrow Z_1} = E(q_0q_1q_2, n_0n_1n_6n_7) \quad (19)$$

$$E_{Z_1 \rightarrow Z_2} = E(q_0q_1q_2, n_2n_3n_4n_5) \quad (20)$$

$$E_{Z_2 \rightarrow Z'_3} = E(q_3q_4q_5, n_3n_4n_5) \quad (21)$$

$$E_{Z_3 \rightarrow Z_4} = E(q_3q_4q_5, n_0n_1n_2) \quad (22)$$

$$E_{Z_4 \rightarrow Y'} = E(q_6q_7q_8, n_0n_1n_2) \quad (23)$$

The register assignments for the E -arrays are:

$$[E_{X \rightarrow Z_1}] \quad b \leftrightarrow \text{ReIm}_{\text{in}} \quad t_0t_1t_2t_3t_4 \leftrightarrow n_7n_6q_0q_1q_2 \quad r \leftrightarrow n_0n_1\text{ReIm}_{\text{out}} \quad (24)$$

$$[E_{Z_1 \rightarrow Z_2}] \quad b \leftrightarrow n_5 \quad t_0t_1t_2t_3t_4 \leftrightarrow n_4n_3q_0q_1q_2 \quad r \leftrightarrow n_2\text{ReIm} \quad (25)$$

$$[E_{Z_2 \rightarrow Z'_3}] \quad b \leftrightarrow n_5 \quad t_0t_1t_2t_3t_4 \leftrightarrow n_4n_3q_3q_4q_5 \quad r \leftrightarrow \text{ReIm} \quad (26)$$

$$[E_{Z_3 \rightarrow Z_4}] \quad b \leftrightarrow n_2 \quad t_0t_1t_2t_3t_4 \leftrightarrow n_1n_0q_3q_4q_5 \quad r \leftrightarrow \text{ReIm} \quad (27)$$

$$[E_{Z_4 \rightarrow Y'}] \quad b \leftrightarrow n_2 \quad t_0t_1t_2t_3t_4 \leftrightarrow n_1n_0q_6q_7q_8 \quad r \leftrightarrow \text{ReIm} \quad (28)$$

I leave it to you to verify that the register assignments in Eqs. (13)–(28) are **consistent**, in the sense that no additional transposing is needed to do the FMAs and MMAs listed at the end of §1.1. (Achieving this consistency is the most difficult part of designing the FFT microkernel. I think I got it right, but let me know if something seems wrong!)

From the E -array register assignments in Eqs. (24)–(28), it may appear that $8 + 4 + 2 + 2 + 2 = 18$ registers are needed to store the E -arrays. However, the register footprint may be reduced as follows:

- For the $E_{X \rightarrow Z_1}$ array (8 registers/thread), it’s possible to store the array in the following “packed” form (4 registers/thread):

$$[E_{X \rightarrow Z_1}] \quad b \leftrightarrow \text{ReIm} \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow n_7 n_6 q_0 q_1 q_2 \quad r \leftrightarrow n_0 n_1 \quad (29)$$

and “unpack” on the fly with a few `__byte_perm` instructions and sign flips. (Let me know if this description is too cryptic and I can give more detail!)

- The $E_{Z_3 \rightarrow Z_4}$ array (2 registers/thread) has the same register content as the $n_2 = 0$ subset of the $E_{Z_1 \rightarrow Z_2}$ array (4 registers/thread). Therefore, no additional registers are needed to store $E_{Z_3 \rightarrow Z_4}$.
- The $E_{Z_4 \rightarrow Y'}$ array (2 registers/thread) has the same register content as the $E_{Z_2 \rightarrow Z'_3}$ array (2 registers/thread). Therefore, no additional registers are needed to store $E_{Z_4 \rightarrow Y'}$.

Using these observations, the E -array footprint can be reduced to $4 + 4 + 2 = 10$ registers/thread.

1.3 Flop-counting

The FFT microkernel uses:

- 8 `m16n8k8` MMAs (2 SM-cycles each).
- 16 `m16n8k16` MMAs (4 SM-cycles each).
- 64 `__half2` FMAs (0.5 SM-cycles each).
- 16 `__shfl_sync` instructions (2 SM-cycles each).
- 32 `__byte_perm` instructions (0.5 SM-cycles each).

The total cost is:

$$(8 \cdot 2) + (16 \cdot 4) + (64 \cdot 0.5) + (16 \cdot 2) + (32 \cdot 0.5) = 160 \frac{\text{SM-cycles}}{\text{FFT}} \quad (30)$$

The real-time FFT throughput (per GPU) is:

$$\frac{(4 \text{ cylinders}) \cdot (2 \text{ polarizations}) \cdot (16 \text{ freqs})}{2.56 \mu\text{s}} = 5 \times 10^7 \frac{\text{FFTs}}{\text{sec}} \quad (31)$$

The L40S compute envelope is:

$$(142 \text{ SMs}) \cdot (2100 \text{ MHz}) = 3 \times 10^{11} \frac{\text{SM-cycles}}{\text{sec}} \quad (32)$$

Combining these estimates, the “GPU load” (defined as fraction of GPU compute needed to run at real-time) is:

$$\text{GPU load} = \frac{(160 \text{ SM-cycles/FFT}) \cdot (5 \times 10^7 \text{ FFTs/sec})}{(3 \times 10^{11} \text{ SM-cycles/sec})} = 0.0267 \quad (33)$$

That is, we predict that **2.67% of the GPU is needed**. (This is not the cost of the entire FRB beamforming kernel – just the first FFT!)

A Local/warp transposes

This appendix is cut-and-paste from the CHORD FRB beamformer doc, but included to make these notes self-contained.

A.1 Local transpose operation

Suppose we have a situation where each thread holds two registers, and each register stores four 8-bit quantities. In our register assignment notation, we write:

$$b_1b_0 \leftrightarrow XY \quad r \leftrightarrow Z \quad (34)$$

to indicate that the three “physical” index bits b_1b_0r correspond to “logical” index bits XYZ , where the meaning of the logical bits depends on the larger context. (We have omitted the physical thread index bits $t_4t_3t_2t_1t_0$, since the operation we will describe is thread-local.)

Now suppose that we want to change the register assignment, by swapping the roles of physical index bits b_0 and r , to get the register assignment:

$$b_1b_0 \leftrightarrow XZ \quad r \leftrightarrow Y \quad (35)$$

We will call this a “local transpose” operation, since it shuffles data between different registers of the same thread.

Similarly, we might want to transpose physical index bits b_1 and r , so that we obtain the register assignment:

$$b_1b_0 \leftrightarrow ZY \quad r \leftrightarrow X \quad (36)$$

Either of the local transpose operations defined in Eqs. (35), (36) can be implemented with two calls to the `__byte_perm()` cuda intrinsic. According to my benchmark, `__byte_perm()` has a throughput of two instructions per cycle (i.e. one local transpose per cycle).

A.2 Warp transpose operation

Now suppose we have a situation where each thread in a warp holds two 32-bit registers:

$$r \leftrightarrow X \quad t_4t_3t_2t_1t_0 \leftrightarrow Y_4Y_3Y_2Y_1Y_0 \quad (37)$$

where we are now keeping track of the 5-bit thread index $t = [t_4t_3t_2t_1t_0]_2$, but not keeping track of byte index bits (i.e. we are treating register contents as 32-bit, not 4×8 -bit).

Suppose that we want to transpose index bits r and t_i , so that we obtain the register assignment:

$$r \leftrightarrow Y_i \quad t_4t_3t_2t_1t_0 \leftrightarrow Y_4 \cdots \underbrace{X}_{\text{replacing } Y_i} \cdots Y_0 \quad (38)$$

This can be done efficiently with one warp shuffle instruction as follows:¹

```
int in0 = ...; // contents of register 0
int in1 = ...; // contents of register 1

int i = ...; // thread index bit 0 <= i < 5
int bit = 1 << i;
```

¹Based on my microbenchmarks, the code below will be warp shuffle limited, i.e. the computation of `bit/flag` and the conditional assignments involving `src/dst` are faster than the warp shuffle and can run in parallel. I also find that warp shuffle throughput is 16 shuffles per clock cycle (where a warp shuffle involving all 32 threads in a warp is defined as 32 shuffles). A puzzle here is that this contradicts nvidia’s throughput tables at <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#maximize-instruction-throughput>, which claim 32 shuffles per cycle. If you have any insight on how to get 32 shuffles per cycle, that would be really valuable, since the FRB *search* kernels are sometimes warp shuffle bound. (For the beamforming kernel which is the subject of this note, the cost of warp shuffles turns out to be small (Table ??), but I thought the larger issue was worth mentioning.

```

bool flag = (threadIdx.x & bit) != 0;

int src = flag ? in0 : in1;
int dst = __shfl_xor_sync(0xffffffff, src, bit);

// Compiles to conditional move, not warp-divergent branch.
(flag ? out0 : out1) = dst;

```

We will call this a “warp transpose” operation, since it shuffles data between different threads in the same warp.

B Float16 tensor core reference

This appendix is cut-and-paste from the CHORD FRB beamformer doc, but included to make these notes self-contained.

B.1 Float16 m16n8k8

The PTX instruction `mma.sync.aligned.m16n8k8.row.col.f16.f16.f16.f16` performs the following matrix multiplication $C = AB$:

$$[(16 \times 8) \text{ float16 } A_{ij}] \quad b_0 \leftrightarrow j_0 \quad r_0 \leftrightarrow i_3 \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow j_1 j_2 i_0 i_1 i_2 \quad (39)$$

$$[(8 \times 8) \text{ float16 } B_{jk}] \quad b_0 \leftrightarrow j_0 \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow j_1 j_2 k_0 k_1 k_2 \quad (40)$$

$$[(16 \times 8) \text{ float16 } C_{ik}] \quad b_0 \leftrightarrow k_0 \quad r_0 \leftrightarrow i_3 \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow k_1 k_2 i_0 i_1 i_2 \quad (41)$$

This instruction performs 2048 flops, and costs 2 SM-cycles on an A40.

B.2 Float16 m16n8k16

The PTX instruction `mma.sync.aligned.m16n8k16.row.col.f16.f16.f16.f16` performs the following matrix multiplication $C = AB$:

$$[(16 \times 16) \text{ float16 } A_{ij}] \quad b_0 \leftrightarrow j_0 \quad r_0 r_1 \leftrightarrow i_3 j_3 \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow j_1 j_2 i_0 i_1 i_2 \quad (42)$$

$$[(16 \times 8) \text{ float16 } B_{jk}] \quad b_0 \leftrightarrow j_0 \quad r_0 \leftrightarrow j_3 \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow j_1 j_2 k_0 k_1 k_2 \quad (43)$$

$$[(16 \times 8) \text{ float16 } C_{ik}] \quad b_0 \leftrightarrow k_0 \quad r_0 \leftrightarrow i_3 \quad t_0 t_1 t_2 t_3 t_4 \leftrightarrow k_1 k_2 i_0 i_1 i_2 \quad (44)$$

This instruction performs 4096 flops, and costs 4 SM-cycles on an A40.